

Design Decisions, Implementation Notes & Answers to All Assignment Questions

This section explains **why each architectural choice was made** and provides explicit **answers to all required “suggest and implement” tasks** from the DevOps challenge.

1. Kubernetes Cluster (Kind)

Reasoning

- Local, lightweight Kubernetes cluster.
- Supports multi-node topology.
- No cloud accounts required.
- Works perfectly with PVCs, StatefulSets, CronJobs, and NetworkPolicies.

Implementation

A 3-node cluster is created using `kind-config.yaml`:

- 1 control-plane node
 - 2 worker nodes
 - Host port `30080` mapped for browser-based testing
-

2. Database Cluster — MariaDB StatefulSet (Persistent)

Reasoning

- StatefulSets provide stable DNS names (`mysql-0`, `mysql-1`).
- Each replica gets its own PVC for durability.
- Ideal for database consistency and crash recovery.

Implementation

- MariaDB deployed as a StatefulSet.
 - VolumeClaimTemplates ensure persistent storage.
 - A headless service provides stable DNS for replicas.
-



3. Web Server (Nginx) With Multi-Replica Deployment

Reasoning

- Nginx is fast, small, and stable.
- Great for demonstrating configuration injection and init containers.

Implementation

- Deployment with multiple replicas.
- Custom Nginx config mounted via ConfigMap.
- Init container modifies HTML before Nginx starts.
- HTML page displays:
 - Pod IP
 - `serving-host = Host-xxxxx` (last 5 chars of pod name)

Example:

```
web-server-7f89cf47bf-25gxj → Host-5gxj
```

✓ 4. Restrict DB Access — Only Web Pods Can Reach MySQL

Assignment Requirement

“Suggest and implement a way to only allow the web server pods to initiate connections to the database pods on port 3306.”

Choice Made: Kubernetes NetworkPolicy

Reasoning

- Native Kubernetes security control.
- Blocks all east-west traffic except what you explicitly allow.
- No service mesh required.

Implementation

NetworkPolicy allows:

```
from:  
  - podSelector:  
    matchLabels:  
      app: devops-demo-web  
ports:  
  - port: 3306
```

All other traffic → **denied**.

✓ 5. Disaster Recovery for Database

Assignment Requirement

“Suggest and implement a disaster recovery solution for the DB.”

Choice Made: mysqldump CronJob → backup PVC

Reasoning

- CronJobs automate recurring backups.
- `mysqldump` is universally compatible.
- Backup PVC stores dumps separate from DB storage.

Implementation

CronJob:

- Runs `mysqldump -A` daily (or more frequently).
 - Dumps stored in `/var/lib/mysql-backups` on a PVC.
 - Supports manual execution.
-



6. Flexible Way to Attach Pods to a Secondary Network

Assignment Requirement

“Find and implement a flexible way to connect the Pod to a new network other than the Pods network with proper routes. No LoadBalancer needed.”

Choice Made: Multus CNI secondary network

Reasoning

- Multus allows pods to attach to **multiple network interfaces**.

- Industry-standard for multi-NIC Kubernetes workloads.
- Clean, declarative, and optional on a per-pod basis.

Implementation

Pods add:

```
annotations:  
  k8s.v1.cni.cncf.io/networks: macvlan-conf
```

This attaches the pod to an additional network with its own routing table.



7. Schedule Specific DB Replicas on Specific Nodes

Assignment Requirement

“Schedule specific replicas of the database cluster on specific K8s nodes.”

Choice Made: Node labels + NodeAffinity

Reasoning

- StatefulSets ensure stable replica numbering (`mysql-0`, `mysql-1`).
- NodeAffinity allows deterministic scheduling.
- Common in HA DB architectures (one replica per node).

Implementation

Nodes labeled:

```
kubectl label node worker-1 db=node-1  
kubectl label node worker-2 db=node-2
```

StatefulSet includes:

```
nodeAffinity:  
  requiredDuringSchedulingIgnoredDuringExecution:  
    nodeSelectorTerms:  
      - matchExpressions:  
          - key: db  
            operator: In  
            values:  
              - node-1
```

This pins DB replicas to specific nodes.

8. Golang Pod Monitoring Application

Reasoning

- Uses the official Kubernetes client-go library.
- Informers provide **real-time**, event-driven updates.
- Required events:
 - Pod created
 - Pod updated
 - Pod deleted

Implementation

- Watches all pods in all namespaces.
- Prints structured logs:
 - Timestamp
 - Event type
 - Pod namespace/name

- Pod IP (if available)
- Phase changes for update events

Example:

```
2025-11-23T20:22:29Z - ADDED: web-75546bd9dd-9t4bm (IP:10.244.2.14)
2025-11-23T20:40:23Z - UPDATED: web-75546bd9dd-vtsd4 (phase=Failed)
2025-11-23T20:39:54Z - DELETED: web-75546bd9dd-hgnwg
```

Architecture Diagram

Placed inside the repository as `architecture.png`.